

TASKIT OS: AI-Native Agency Operating System Strategy

Operational infrastructure, AI systems, automation, financial intelligence, and enterprise platform direction

North Star

TASKIT OS should become the AI-powered operational infrastructure layer for modern agencies and service businesses. The winning platform is not a productivity app. It is the system that connects client demand, scope, capacity, production, approval, delivery evidence, invoice, payment, profitability, and executive decisions.

Executive Thesis

TASKIT has the right early foundations: workspace tenancy, clients, campaigns, briefs, deliverables, tasks, approvals, invoices, audit logs, jobs, search indexes, analytics metrics, and AI memory. The next strategic step is to unify these foundations into one governed Agency Execution Graph.

Generic platforms capture tasks. TASKIT must capture agency value creation. Every client promise should become structured scope, every deliverable should become auditable value, every approval should become invoice readiness, every invoice should connect to work evidence, and every operational signal should become AI-assisted intelligence.

What TASKIT Must Become

SMARTER THAN CLICKUP

AI that understands agency delivery economics, not generic task text.

MORE OPERATIONAL THAN MONDAY.COM

Native state machines, capacity, delivery-to-cash, and governance.

MORE AI-NATIVE THAN ASANA

Role-specific agents embedded in agency workflows with financial context.

MORE INTEGRATED THAN NOTION

Knowledge, execution, approvals, files, billing, and automation in one graph.

MORE AGENCY-FOCUSED THAN ALL

Deep creative approvals, client portals, retainers, margin, and service operations.

HARDER TO REPLACE

Client adoption, automation history, operating memory, and financial records create lock-in.

Platform Analysis Lenses

AREA	DIRECTION
Operational infrastructure	Owns the agency execution graph from request intake through delivery, approval, invoice, and payment.
AI operating system	Permission-scoped copilots and agents reason over live operations, propose actions, execute safely, and learn from outcomes.
Business intelligence layer	Turns tasks, approvals, invoices, activity, time, and client signals into executive-grade operating metrics.
Workflow execution engine	Moves work through governed state machines, automations, retries, escalations, and audit trails.
Financial operating layer	Connects scope, cost, time, delivery evidence, invoices, collections, and profitability.
Automation platform	Lets agencies encode repeatable operating models as versioned rules, workflows, and agentic playbooks.
Collaboration platform	Unifies internal teams, clients, reviewers, finance, and AI agents in shared operating rooms.
Enterprise governance system	Controls permissions, audit, approvals, agent tools, portal access, policies, and compliance exports.
Knowledge system	Builds a searchable graph of briefs, SOPs, files, comments, reports, decisions, and client history.
Decision-making engine	Predicts risk, recommends actions, explains tradeoffs, and converts decisions into workflow changes.

Current Weaknesses and Missing Systems

AREA	DIRECTION
Execution model	Project still acts as a legacy campaign backing model, so product language and database truth remain partially split.
Financial depth	No contracts, SOWs, retainers, cost rates, expenses, time entries, change orders, or revenue recognition primitives yet.
Capacity intelligence	No skills matrix, availability, allocations, PTO, department load, utilization targets, or forecasted demand model.
AI runtime	Current AI is useful but mostly deterministic retrieval plus limited actions, not a governed multi-agent operating layer.
Automation infrastructure	JobRun exists, but TASKIT needs workflow definitions, rule versions, simulations, idempotency, retries, and run observability.
Search and knowledge	SearchIndex is keyword-first; the platform needs hybrid search, entity graph traversal, embeddings, and source citations.
Permissions	The current role system should evolve into custom roles, ABAC, client-scoped access, field controls, and AI tool permissions.
Analytics	AnalyticsMetric is a good foundation, but TASKIT needs a semantic warehouse for margin, risk, capacity, SLA, and benchmarks.
Collaboration	The client portal should become a branded operating room for approvals, requests, files, invoices, reports, and decisions.

Retention loops	The strongest future loops are client portal adoption, automation dependency, historical intelligence, and financial records.
Network effects	Anonymous agency benchmarks, workflow templates, partner playbooks, and client-side adoption can create defensibility.
Monetization	Expansion should come from AI seats, agent runs, automation volume, client portals, finance, storage, governance, search, and API usage.

Competitor Gaps TASKIT Can Dominate

AREA	DIRECTION
ClickUp	Strong broad work management, Brain, connected search, agents, and MCP. Still horizontal and not deeply tied to agency margin, creative approvals, or delivery-to-cash.
monday.com	Excellent configurable workflows, dashboards, automations, and enterprise messaging. It does not natively encode agency delivery economics or client approval governance.
Asana	Strong Work Graph, AI Teammates, workflow AI, and enterprise controls. Its advantage is broad team coordination, not vertical agency operating intelligence.
Notion	Very strong knowledge, docs, and enterprise search. It is weaker as a governed execution engine for delivery, finance, and operations.
ServiceNow / Salesforce	Enterprise agent governance and workflow maturity are the bar. TASKIT should copy the control-plane discipline but specialize in agencies.
Rippling	Proves that permissions plus workflows across business data create deep dependency. TASKIT should apply that model to agency operations.

The market direction is clear: work platforms are adding enterprise search, AI agents, agent governance, connected apps, and workflow automation. TASKIT should not try to out-horizontal the horizontal platforms. The advantage is to go deeper into agency operations than they can justify.

Core Platform Redesign

The canonical platform object should be the Agency Execution Graph. Recommended graph nodes include Workspace, Business Unit, Client, Contact, Contract, SOW, Retainer, Service Catalog Item, Intake Request, Campaign, Brief, Deliverable, Asset Version, Review Room, Approval Decision, Task, Subtask, Dependency, Time Entry, Cost Rate, Expense, Invoice Line, Payment, Report, Decision, Risk, Automation Run, AI Agent Run, and Knowledge Chunk.

Primary lifecycle

Client demand to scope to capacity plan to brief to deliverables to approvals to delivery evidence to invoice to payment to profitability to executive decision to automated next action.

- Every important record should emit events into an immutable Operational Ledger.
- Every event should be usable by dashboards, automations, AI agents, audit logs, search, and predictions.
- SQL should remain the source of truth for measurable business facts.
- Vector retrieval should be used for semantic content such as briefs, comments, files, SOPs, contracts, and reports.
- Graph traversal should connect clients, work, decisions, approvals, invoices, people, and risk signals.

Future Architecture Direction

System of record

PostgreSQL and Prisma remain authoritative for transactional entities and tenant isolation.

Event spine

Domain events flow into activity, audit, search, analytics, automations, and AI evaluation.

Workflow engine

Versioned workflow definitions govern state transitions, SLAs, retries, and escalations.

AI control plane

Tool registry, policy engine, agent runs, approvals, memory, and evaluations.

Knowledge plane

Hybrid search, embeddings, entity graph, citations, connector ingestion, and permission filters.

Intelligence warehouse

Semantic metrics for risk, margin, capacity, client health, SLA, utilization, and benchmarks.

Strategic Initiatives

Agency Execution Graph

Why it matters	The product must model the real lifecycle of agency value, not just projects and tasks.
Competitive advantage	Generic competitors cannot easily replicate a vertical graph spanning scope, creative review, approval, cash, margin, and client health.
Retention	Historical operations, approvals, assets, invoices, and decisions compound into the customer workspace.
Customer dependency	Teams run daily decisions from the graph, making migration operationally painful.
Platform value	Creates the system of record and system of intelligence for the agency.
Complexity	High. Requires schema evolution, event discipline, graph relationships, and migration from project-language compatibility.
Monetization	Core premium packaging, enterprise plans, implementation services, and API expansion.
AI opportunities	AI can reason over connected entities and answer why work is late, unprofitable, blocked, or ready to invoice.
Automation opportunities	Every graph transition can trigger workflows, notifications, risk updates, billing actions, and reports.
Enterprise value	Supports audit, governance, multi-workspace rollups, and portfolio operating models.
Long-term scale	Becomes stronger as every workspace event improves graph density, benchmarks, and models.

Delivery-to-Cash Engine

Why it matters	Agencies care about cash conversion, not only work completion.
Competitive advantage	Links approved deliverables to invoice readiness, payment reminders, collections, and revenue exposure.
Retention	Finance teams become daily users and historical billing records anchor the platform.
Customer dependency	Replacing TASKIT means rebuilding how delivery evidence maps to cash.
Platform value	Turns operations into measurable financial outcomes.
Complexity	High. Requires invoice line linkage, payment states, contract terms, accounting integrations, and revenue events.
Monetization	Finance module, payment processing, recurring billing, collections automation, accounting connectors.
AI opportunities	Finance Agent drafts invoices, explains revenue risk, predicts late payment, and recommends collections paths.
Automation opportunities	Approved deliverable creates invoice draft, overdue invoice triggers reminders and account-owner escalation.
Enterprise value	CFO dashboards, revenue controls, audit trails, approval-to-invoice evidence packages.
Long-term scale	Extends to retainers, subscriptions, usage billing, revenue recognition, and multi-currency operations.

Strategic Initiatives Continued

Profitability Intelligence

Why it matters	Agency growth breaks when teams cannot see which clients, services, or projects destroy margin.
Competitive advantage	Generic PM tools rarely connect work, cost, time, rates, scope changes, and collections.
Retention	Owners and finance depend on TASKIT for margin decisions and pricing strategy.
Customer dependency	The more history captured, the harder it is to rebuild true client and service profitability elsewhere.
Platform value	Moves TASKIT from productivity tool to executive operating system.
Complexity	High. Requires time tracking, billable status, cost rates, service catalog, expenses, and clean invoice mapping.
Monetization	Advanced analytics tier, CFO cockpit, margin forecasting, pricing intelligence.
AI opportunities	AI predicts margin erosion, detects scope creep, and suggests pricing or staffing changes.
Automation opportunities	Trigger margin alerts, change-order drafts, budget warnings, and invoice adjustment workflows.
Enterprise value	Board reporting, department P&L, client portfolio analysis, finance governance.
Long-term scale	Enables benchmarking by service line, market, client segment, and team structure.

Capacity Forecasting System

Why it matters	Agencies need to know if promised work can be delivered before accepting it.
Competitive advantage	TASKIT can become the COO layer by connecting demand, skills, availability, deadlines, and profitability.
Retention	Leadership forms a weekly planning habit around capacity and risk.
Customer dependency	Staffing, hiring, outsourcing, and prioritization decisions become TASKIT-native.
Platform value	Prevents overcommitment, burnout, missed deadlines, and margin loss.
Complexity	High. Requires availability, skills, allocations, PTO, departments, utilization targets, and forecast models.
Monetization	Resource planning add-on, executive seats, forecasting tier.
AI opportunities	Resourcing Agent recommends rebalancing, hiring, contractor needs, and delivery tradeoffs.
Automation opportunities	Auto-escalate overload, suggest reassignment, block new commitments, and update timelines.
Enterprise value	Multi-team workforce planning and utilization governance.
Long-term scale	Can forecast across workspaces, business units, services, and geographies.

Governed Multi-Agent Runtime

Why it matters	The market has moved from AI chat to agents that safely act inside workflows.
Competitive advantage	TASKIT can offer agency-specialized agents with better context and safer controls than generic assistants.
Retention	Agents become embedded operational labor that customers rely on every day.
Customer dependency	Customers build internal workflows around agent outputs, approvals, and histories.
Platform value	Turns TASKIT into an AI operating system rather than a dashboard.
Complexity	Very high. Requires agent tables, tool registry, approval gates, policy engine, run logs, evals, and rollback paths.
Monetization	AI seats, agent run volume, premium agent packs, enterprise governance.
AI opportunities	Executive, producer, finance, account, resource, QA, comms, automation, and knowledge agents.
Automation opportunities	Agents can draft, route, escalate, update records, generate reports, and trigger workflows.
Enterprise value	Agent observability, least-privilege tools, approval policies, audit logs, usage analytics.
Long-term scale	Agents should be modular and share the same action, policy, and event substrate.

Strategic Initiatives Continued

Automation Orchestration Layer

Why it matters	Sticky SaaS products let customers encode their operating model into the platform.
Competitive advantage	TASKIT can automate agency-specific workflows competitors treat as manual coordination.
Retention	Every automation reduces labor and increases switching cost.
Customer dependency	Processes become invisible infrastructure customers do not want to rebuild.
Platform value	Drives measurable time savings and operational consistency.
Complexity	High. Needs rule DSL, workflow definitions, versions, simulations, queue observability, idempotency, and dead-letter recovery.
Monetization	Automation limits, run volume, workflow packs, enterprise queues, partner templates.
AI opportunities	Natural-language automation builder compiles into validated drafts rather than executing blindly.
Automation opportunities	Trigger-condition-action workflows across delivery, approvals, finance, notifications, and reports.
Enterprise value	Versioning, approvals, simulations, run logs, retries, and policy checks.
Long-term scale	Can evolve from monolith workers to dedicated queue, workflow, and orchestration services.

Client Operating Room

Why it matters	External client participation creates a stickiness generic internal tools cannot match.
Competitive advantage	Branded portals turn TASKIT into part of the agency service experience.
Retention	Clients approve, request, comment, pay, and consume reports inside TASKIT.
Customer dependency	Leaving TASKIT disrupts the agency-client operating rhythm.
Platform value	Reduces account-management overhead and increases client trust.
Complexity	Medium to high. Requires client-safe permissions, tokens, branded portals, custom domains, data rooms, and audit.
Monetization	Client portals, guest seats, white label, custom domains, client report packs, storage.
AI opportunities	Client Comms Agent drafts updates, approval reminders, meeting recaps, and client-safe reports.
Automation opportunities	Request intake creates briefs, approvals trigger delivery, payment reminders route through portal.
Enterprise value	Client-scoped audit trails, expiring access, data-room controls, legal export.
Long-term scale	Can support multi-client workspaces, enterprise stakeholders, and agency resale models.

Enterprise Knowledge Graph and Search

Why it matters	Agencies lose time because knowledge is scattered across briefs, files, comments, docs, emails, and chats.
Competitive advantage	TASKIT can make every client, decision, deliverable, SOP, and invoice searchable with permissioned citations.
Retention	Accumulated knowledge becomes a durable memory layer.
Customer dependency	Teams use TASKIT as the fastest path to understanding any account or campaign.
Platform value	Reduces context switching and improves decision quality.
Complexity	High. Needs hybrid search, embeddings, source connectors, chunking, permissions, and graph traversal.
Monetization	Enterprise search tier, connector pricing, storage, AI knowledge packs.
AI opportunities	RAG with SQL-first facts, vector semantic retrieval, graph context, citations, and memory writes.
Automation opportunities	Index every event, file, brief, comment, SOP, report, and external connector update.
Enterprise value	Permission inheritance, source controls, citation audit, data retention, zero-training posture.
Long-term scale	Search layer can become a shared service used by AI, command palette, reports, and agents.

Strategic Initiatives Continued

Predictive Delivery Risk

Why it matters	The best operations platform does not only report problems; it predicts them early.
Competitive advantage	Agency-specific risk models can use approval age, client response time, revision count, workload, deadline compression, and invoice exposure.
Retention	Leaders check TASKIT daily because it shows what will break next.
Customer dependency	Risk histories and recommended recovery actions become operational memory.
Platform value	Improves delivery reliability, client trust, and margin protection.
Complexity	Medium to high. Starts rule-based, then trains models as historical data grows.
Monetization	Predictive intelligence tier, executive dashboards, SLA reporting.
AI opportunities	Risk Agent explains causes, confidence, recommended next action, and likely outcome.
Automation opportunities	Auto-create recovery plans, escalate blockers, reschedule work, and notify stakeholders.
Enterprise value	Risk scoring governance, explainability, model evaluation, and auditability.
Long-term scale	Improves as the platform captures more operational outcomes and benchmarks.

Scope and Change Control

Why it matters	Untracked client changes are one of the largest sources of agency margin leakage.
Competitive advantage	TASKIT can detect, formalize, approve, and bill change requests inside the same delivery graph.
Retention	Agencies depend on TASKIT to protect margins and document client accountability.
Customer dependency	Contracts, approvals, revisions, and change orders become interconnected records.
Platform value	Turns ambiguous work into governed commercial decisions.
Complexity	Medium. Requires SOWs, scope items, change requests, approval routing, and invoice linkage.
Monetization	Finance ops, contract module, change-order automation.
AI opportunities	AI detects scope creep in comments, requests, and briefs, then drafts change orders.
Automation opportunities	Out-of-scope request triggers manager review, client approval, and invoice item draft.
Enterprise value	Contract-aware approvals, legal evidence, client dispute protection.
Long-term scale	Extends into service catalogs, packaged offerings, and retainer governance.

Executive Command Center

Why it matters	The platform must become the first tab agency leaders open each morning.
Competitive advantage	Combines operations, finance, client health, capacity, approvals, and AI recommendations in one operating cockpit.
Retention	Creates daily executive dependency and top-down renewal pressure.
Customer dependency	Leadership workflows, reports, and decisions become TASKIT-native.
Platform value	Makes TASKIT board-level infrastructure.
Complexity	Medium. Needs semantic metrics, role-based dashboards, snapshots, drilldowns, and AI briefs.
Monetization	Executive seats, advanced analytics, scheduled board/client reports.
AI opportunities	Executive Chief of Staff creates daily briefs, weekly reviews, and decision memos.
Automation opportunities	Scheduled reports, threshold alerts, portfolio health updates, and decision follow-through.
Enterprise value	Multi-workspace rollups, board packs, audit-backed metrics.
Long-term scale	Can support agency groups, departments, locations, and portfolio companies.

Marketplace, Benchmarks, and Playbooks

Why it matters	A platform becomes more defensible when users and partners add reusable operating knowledge.
Competitive advantage	Vertical playbooks and anonymous benchmarks are hard for horizontal platforms to build quickly.
Retention	Customers keep returning for workflow improvements, agency standards, and performance comparisons.
Customer dependency	Templates, automations, reports, and AI memories become tuned to each agency.
Platform value	Creates ecosystem and data-network effects.
Complexity	Medium. Requires template schema, marketplace governance, anonymization, partner APIs, and quality controls.
Monetization	Paid templates, partner revenue share, benchmark reports, certification programs.
AI opportunities	AI recommends playbooks based on agency type, risk pattern, and maturity.
Automation opportunities	Installable workflow packs for onboarding, approvals, retainers, collections, and reporting.
Enterprise value	Approved template libraries, centralized governance, compliance-ready playbooks.
Long-term scale	Expands the product without every workflow being built by the core team.

AI-First Architecture

Prompt to Intent Router to Policy Engine to SQL and Graph Retrieval to Vector Retrieval to Tool Planner to Human Approval Gate to Action Executor to Event Ledger to Memory Update to Evaluation Log.

Tool registry

Defines exactly what each agent can read, draft, update, send, approve, or execute.

Policy engine

Checks role, workspace, client, entity, field, and action permissions before retrieval or action.

Agent run log

Stores every plan, tool call, input, output, error, citation, approval, and final action.

Memory store

Separates user memory, workspace memory, client memory, workflow memory, and decision memory.

Hybrid retrieval

SQL for facts, vector search for knowledge, graph context for relationships, citations for trust.

Evaluation harness

Tests hallucination, permission leaks, action accuracy, regression, and agent business impact.

Multi-Agent Architecture

AREA	DIRECTION
Executive Chief of Staff	Daily operating brief, bottlenecks, decisions, revenue risk, and follow-through.
Account Director Agent	Client health, stakeholder follow-ups, renewal risk, expansion signals, and communication drafts.
Producer Agent	Delivery risk, dependencies, approval queues, timeline recovery, and production priorities.
Finance Agent	Invoice readiness, collections, profitability, revenue exposure, and payment workflow recommendations.
Resourcing Agent	Capacity forecast, overload detection, skills matching, staffing needs, and reassignment plans.
Creative QA Agent	Compares deliverables against brief, brand rules, specifications, and client feedback.
Client Comms Agent	Drafts status updates, approval reminders, recap emails, and client-safe executive reports.
Automation Architect	Turns natural language into validated automation drafts with trigger, condition, action, and approval policy.
Knowledge Librarian	Maintains SOPs, decisions, briefs, reusable learnings, and knowledge graph hygiene.

AI Memory Systems

- Short-term memory: recent conversation and active workspace context.
- User memory: preferences, reporting format, focus areas, tone, and personal workflow habits.
- Workspace memory: SOPs, policies, recurring bottlenecks, preferred metrics, and approved templates.
- Client memory: stakeholders, preferences, risks, renewal context, approval behavior, and communication history.
- Decision memory: what leadership decided, why, what data supported it, and which workflow changed after the decision.

- Memory writes should be explicit, permission-scoped, auditable, and reversible.

Automation Infrastructure

- Create a workflow DSL with triggers, conditions, actions, wait states, approvals, retries, and compensating actions.
- Version every automation so changes can be reviewed, tested, rolled back, and audited.
- Add simulation mode so owners can see what a rule would have done over historical events before activation.
- Require human approval for external communication, financial action, destructive update, or high-risk AI-generated change.
- Use idempotency keys and action ledgers so retries do not duplicate invoices, comments, reminders, or tasks.
- Expose queue observability for active, failed, delayed, completed, and dead-letter workflow runs.

Financial Intelligence Systems

Contracts and SOWs Define scope, terms, rates, deliverables, approval requirements, and billing schedule.	Retainers Track recurring commitments, included hours or deliverables, overages, and renewal risk.
Time and cost Capture billable time, non-billable time, internal cost rates, freelancers, and expenses.	Change orders Detect out-of-scope requests, route approvals, and convert approved changes into invoice lines.
Revenue recognition Map delivery milestones, approvals, invoice status, payment state, and deferred revenue.	Profitability dashboard Margin by client, service, campaign, deliverable, team, manager, and month.

Executive Dashboards and Predictive Systems

- Executive daily brief: what changed, what is at risk, what needs a decision, and where cash is exposed.
- Delivery risk: late tasks, approval age, revision count, dependency blockers, workload pressure, and client response delays.
- Profitability: budget burn, margin forecast, unbilled work, change-order leakage, and retainer overrun.
- Capacity: utilization, available hours, skills gaps, overload, hiring need, and contractor recommendations.
- Client health: engagement, inactive days, open invoices, approval friction, sentiment, SLA misses, and expansion signals.
- Forecasting: revenue in flight, collections timing, future workload, delivery confidence, and renewal probability.

Enterprise Governance

Identity SAML/SSO, SCIM, two-factor controls, session policy, support access policy.	Permissions Custom roles, ABAC, field-level controls, client-safe policies, agent tool permissions.
Audit Immutable logs for approvals, role changes, AI actions, invoices, portal access, and workflow transitions.	Compliance Data retention, legal hold, export controls, DLP hooks, data residency path, and access reviews.
Governed AI Agent approvals, run logs, evals, source citations, permission checks, and model/provider controls.	Platform admin Multi-workspace governance, policy templates, sandbox workspaces, webhooks, API keys, and admin analytics.

Moats and Lock-In

AREA	DIRECTION
Agency Execution Graph	Hard to replicate because it encodes vertical workflows and financial relationships.
Client portal adoption	External stakeholders use TASKIT for approvals, requests, files, reports, and invoices.
Automation history	Customers depend on invisible workflows that would be expensive to rebuild.
AI memory	Workspace and client-specific context improves over time and makes the assistant more useful.
Delivery risk models	Historical outcomes improve predictions and recovery recommendations.
Approval evidence	Versioned files, comments, decisions, and invoices create legal and operational records.
Benchmarks	Anonymous industry performance data creates network effects across agencies.
Playbook marketplace	Partners and advanced users extend the platform with reusable operating systems.

Monetization and Packaging

- Core OS: workspace, clients, campaigns, briefs, deliverables, tasks, approvals, basic invoices.
- AI Pro: assistant, executive briefs, smart search, report generation, and limited agent actions.
- AI Workforce: role-specific agents, agent runs, approval gates, memory, evaluations, and agent analytics.
- Automation Pro: workflow builder, advanced conditions, run volume, simulations, retries, and observability.
- Finance Ops: retainers, SOWs, change orders, time/cost, profitability, payment workflows, accounting integrations.
- Client Experience: white-label portal, custom domains, client seats, review rooms, storage, branded reports.
- Enterprise: SSO/SCIM, custom roles, audit exports, multi-workspace governance, data controls, API and premium support.

Implementation Roadmap

0-30 days

Freeze canonical hierarchy, define Agency Execution Graph, add time/cost/service catalog schema plan, design automation and agent run tables.

31-60 days

Ship executive metrics from real data, add workflow transitions, start hybrid search indexing, model contracts/SOWs/retainers/change orders.

61-90 days

Launch delivery-to-cash workflows, basic capacity forecasting, agent approval logs, client operating room upgrades, and predictive risk v1.

3-6 months

Add advanced automation builder, AI agents by role, profitability intelligence, governance controls, accounting integrations, and dashboard snapshots.

6-12 months

Introduce marketplace, benchmarks, multi-workspace governance, enterprise search connectors, agent evaluations, and autonomous workflow packs.

Immediate Schema Priorities

- Add Contract, SOW, Retainer, ServiceCatalogItem, ScopeItem, ChangeRequest, TimeEntry, CostRate, Expense, Payment, and RevenueSchedule.
- Add WorkflowDefinition, WorkflowVersion, WorkflowTransition, AutomationRule, AutomationRun, AutomationActionLog, and AutomationApproval.
- Add AiAgent, AiAgentRun, AiToolCall, AiApproval, AiEvaluation, AiKnowledgeChunk, AiMemoryPolicy, and AiSourceCitation.
- Add CapacityProfile, Skill, Availability, Allocation, Department, UtilizationTarget, and ForecastSnapshot.
- Add ClientHealthSnapshot, RiskSignal, RiskModelRun, DecisionRecord, ExecutiveBrief, and BenchmarkMetric.

Final Platform Vision

TASKIT OS can dominate by becoming the platform where every client promise becomes structured work, every deliverable becomes auditable value, every approval becomes revenue readiness, every invoice connects to execution, and every operational signal becomes AI-assisted intelligence.

The long-term advantage is not feature count. It is the agency execution graph, the governed AI control plane, and the financial intelligence layer built from each customer's real operating history.

Sources Used for Market Direction

- ClickUp Brain and agents: <https://help.clickup.com/hc/en-us/articles/12578085238039-What-is-ClickUp-Brain>
- Asana AI Teammates: <https://asana.com/product/ai/ai-teammates>
- monday.com AI work platform: <https://monday.com/>

- Notion Enterprise Search: <https://www.notion.com/product/enterprise-search>
- Salesforce Agentforce platform: <https://www.salesforce.com/platform/agentforce-platform/>
- ServiceNow AI Agent Orchestrator and Workflow Data Fabric:
<https://newsroom.servicenow.com/press-releases/details/2025/ServiceNows-latest-platform-release-adds-to-thousands-of-AI-agents-across-CRM-HR-IT-and-more-for-faster-smarter-workflows-and-maximum-business-impact-03-12-2025-traffic/default.aspx>
- Rippling workflow permissions: <https://www.rippling.com/en-GB/platform/workflows>